
SEMANTIC TOKENIZER FOR ENHANCED LANGUAGE PROCESSING

Sandeep Mehta, Darpan Shah, R. Kulkarni
Inspird, Inc.
Mission Viejo, CA
{smehta, dshah, rkulkarni}@inspird.com

Cornelia Caragea
Computer Science
University of Illinois at Chicago
cornelia@uic.edu

ABSTRACT

Traditionally, NLP performance improvement has been focused on improving models and increasing the number of model parameters. NLP vocabulary construction has remained focused on maximizing the number of words represented through subword regularization. We present a novel tokenizer that uses semantics to drive vocabulary construction. The tokenizer includes a trainer that uses stemming to enhance subword formation. Further optimizations and adaptations are implemented to minimize the number of words that cannot be encoded. The encoder is updated to integrate with the trainer. The tokenizer is implemented as a drop-in replacement for the SentencePiece tokenizer. The new tokenizer more than doubles the number of wordforms represented in the vocabulary. The enhanced vocabulary significantly improves NLP model convergence, and improves quality of word and sentence embeddings. Our experimental results show top performance on two Glue tasks using BERT-base, improving on models more than $50\times$ in size.

1 Introduction

NLP models have two primary components—a deep neural network and a vocabulary of embeddings. Recent improvements in NLP model performance have focused on improving deep networks and increasing model sizes. Interestingly, little attention has been paid to optimizing vocabularies.

An analysis of recent models [1], plotted in Figure 1,¹ shows that model sizes (i.e., number of parameters) have increased by 15,000% over the last few years (excluding 175B parameter GPT-3 [2] or 1T parameters of GPT4). The increase in model size significantly increases training costs. Recent publications show that a single training run for GPT-3 could cost \$12M.² Even BERT-Large [3] training costs reach tens of thousands of dollars. Furthermore, increased model size place additional computational burden during model execution. These costs can have a detrimental effect on NLP innovation.

At the same time, the size of vocabularies has increased only about 100% and the size of embedding vectors has increased about 200%. Hence, over the same period, the fraction of NLP parameters representing the vocabulary has shrunk from 21% in BERT-base [3] to 0.3% in GPT-3 [2].

An average person uses 42,000 root words and hundreds of thousands of wordforms [4]. Technical terminology and jargon add tens of thousands of additional words to the vocabulary. Even some of the largest vocabularies currently in use, DeBERTa [5], have only 128,000 tokens. Hence, NLP vocabularies need to include subwords that can be combined to form multiple words [6]. Words that can not be represented using a single token are segmented into an initial subword followed by as many intermediate subwords as required.

NLP models use a tokenizer to convert strings of characters into a sequence of lexical tokens. Tokenizers also construct the vocabulary of lexical tokens. We present a novel tokenizer that improves NLP performance by improving subword formation and embedding quality through *semantic tokenization*.

¹Credit for figure to Huggingface’s DistilBERT: <https://research.aimultiple.com/gpt/>

²<https://venturebeat.com/2020/06/01/ai-machine-learning-openai-gpt-3-size-isnt-everything/>




Figure 1: Increase in NLP Model Size.

2 Background

Tokenizers consist of three major components:

1. *Trainer*: The trainer builds the NLP vocabulary
2. *Encoder*: Encoders convert strings into a sequence of tokens
3. *Decoder*: Decoders convert token sequences into words and sentences

Commonly used tokenizers include WordPiece [7], Byte-Pair Encoding (BPE) [6] and Unigram [8]. Trainers in these tokenizers build vocabulary through subword regularization. This subword regularization can be seen as an optimization function that maximizes the number of words in a corpus D that can be represented with a given vocabulary size $|V|$. Each tokenizer accomplishes this task with a different approach: BPE forms subwords by maximizing frequency of character sequences, WordPiece maximizes likelihood of subword formation and Unigram minimizes the loss.

Several approaches are available for encoders as well. WordPiece uses a greedy longest-match-first strategy to tokenize a single word: it iteratively picks the longest subword of the remaining text that matches a word in the vocabulary. BPE encoder incrementally finds a set of subwords such that the total number of subwords for encoding the text is minimized. Unigram uses an entropy encoder that minimizes the total subwords for the text.

3 Semantic Tokenizer

Subword regularization approaches have been shown to effectively model out-of-vocabulary words [6]. This behavior is based on the intuition that many words, such as compound words, are formed from smaller common sub-units. Words that cannot be segmented into subwords generate [UNK] tokens which do not have any semantics. However, the subword regularization does not use any semantic information that may be readily available.

An analysis of pretrained WordPiece/BERT [3] and Unigram/ALBERT [9] vocabularies shows that about 75% tokens represent initial subwords and 20% are intermediate subwords. Most intermediate subwords do not have any inherent semantics. Many initial subword tokens are used to model different forms of the same root word. For example, the vocabularies contain eight forms of the word *advise*: *advise*, *advised*, *advisee*, *advisers*, *advises*, *advising*, *advisor*,

and *advisors*. Furthermore, 15 related wordforms such as *advises*, *advisees*, *advisable* are not represented. Semantics quality decline rapidly when words are divided into more than 2 subwords.

We develop a new semantic tokenizer that enhances tokenization using semantics. Subsections below describe a new semantic tokenizer trainer and encoder.

3.1 Trainer for the Semantic Tokenizer

We reformulate subword regularization problem as a dual objective optimization problem:

1. *Maximize the embedding quality*
2. *Maximize the number of words that can be modeled*

To achieve these two objectives, we divide the the vocabulary V into two segments V_1 and V_2 . The segment V_1 generates subwords by focusing on semantics and maximizing embedding quality (Objective 1). The segment V_2 maximizes the number of words that can be modeled and minimizes the number of non-semantic [UNK] tokens (Objective 2). The parameter $f = \frac{|V_1|}{|V|}$ becomes part of the optimization problem and determines the size of the semantic segment of the vocabulary.

Using semantics to guide subword formation can be achieved by focusing on root words of each wordform. Using suffixes as intermediate subwords can improve their semantic efficiency and free up more tokens in the vocabulary. Two linguistics morphological approaches are available for obtaining root words: Stemming and lemmatization. Stems are part of a word common to its inflected variants and lemmas are its canonical form, headword or root word. While lemmas are attractive because they have inherent semantics in modeled language, stems are more attractive for use with subword formation.

For example, forms of the word *advise* are represented with the stem *advis*. Different suffixes such as *##e*, *##er*, *##ing* are used to form wordforms. Since stems are common across wordforms and suffixes are common across the vocabulary, the semantic trainer can represent many wordforms with a smaller set of stems and suffixes. Hence, stemming can better achieve the second tokenizer objective of minimizing number of subwords. Using same stem across wordforms improves semantic similarity between wordforms and improves input embedding quality. Furthermore, stemming improves NLP model training because increased number of occurrences of the stems and suffixes increases masking probability. In the example above, the stem *advis* occurs 230,000 times in the wiki corpus, while the word *advise* only occurs 10,159 times.

The semantic trainer first populates the segment V_1 . The trainer starts by generating a frequency table of the words in the training corpus and processes them in decreasing order of frequency. If stemming is possible for the word, it divides the word into its stem and suffix. If stem or the suffix has been previously been added, the code updates their likelihood. Otherwise, it adds the new stem or suffix to the vocabulary. If the word cannot be stemmed, it is added directly to the vocabulary. Once the first vocabulary segment is filled, the BPE algorithm is called to fill in the second segment V_2 . The BPE algorithm is quite effective at achieving Objective 2 - minimize the number of [UNK] tokens.

We extended SentencePiece [10] with semantic training algorithms and a stemmer. Detailed analysis of vocabularies generated with available stemmers showed that the Snowball (Porter 2) stemmer [11] is the most effective for use in the semantic trainer. The updated code provides a commandline parameter to set the relative size of segments. Experiments with wiki and book corpus showed that allocating 90 - 95% to segment V_1 and 5-10% to segment V_2 produced optimal results.

3.2 Encoder for the Semantic Tokenizer

The encoder for the semantic tokenizer need to ensure usage of correct subwords in tokenization. When multiple subword sequences could represent a word, the encoder needs to select the sequence that uses the stem and suffix over any other non-semantic subword sequence. WordPiece’s greedy longest-match-first strategy is directly applicable to semantic encoding. Hence, WordPiece encoder was used to without modification in the new tokenizer.

4 Experimental Setup

We used BERT-base model using the HuggingFace Transformers library [12] for experiments. Since the focus of our work is to optimize the vocabulary, the benefits demonstrated are likely to scale to all NLP models with varied complexities.

Corpus	Unigram	WordPiece	Semantic
Wiki	20,765	21,506	44,735
Book	19,772	20,655	48,016

Table 1: Number of Wordforms Represented.

Corpus	Unigram		WordPiece		Semantic	
	Ave	CoV	Ave	CoV	Ave	CoV
Wiki	1.9	47%	1.9	47%	1.7	51%
Book	2.0	44%	2.0	47%	2.0	48%

Table 2: Subword Regularization Efficiency.

We used the semantic tokenizer trainer to build a new vocabulary and generated a version compatible with BERT. We generated training data using the new vocabulary, Wiki corpus and Book corpus. We trained the BERT Base model using a train batch size of 256; max sequence length of 256; max predictions per seq of 20; learning rate of $2e-4$ and number of warmup steps = 10,000. The model was trained for 2M steps.

We compared performance of the new semantic tokenizer with pretrained BERT-base/Wordpiece and ALBERT/Unigram vocabularies.

5 Results

5.1 Subword Regularization Efficiency

We compared the new vocabulary against pretrained Wordpiece (BERT) and Unigram (ALBERT) vocabularies. We measured the tokenization efficiency by computing the number of words in Wiki Corpus and Book Corpus that can be represented by 1 or 2 subwords. As can be seen in Table 1, the new vocabulary covers a larger fraction of the words used in each corpus without using generic subwords. The new vocabulary is able to represent 125% more wordforms or use 34% fewer entries to represent the same number of wordforms.

A concern with stemming is that it replaces a word with two subwords, which might increase the number of tokens and reduce the embedding quality. We evaluated number of tokens required to represent the corpus vocabulary and the variation in tokenization. As shown in Table 2, this analysis showed that the semantic trainer actually reduced the average number of subwords per word.

5.2 Embedding Quality

We evaluated the overall performance of the semantic tokenizer against Glue tasks (Table 3). The new tokenizer helped BERT-base outperform the leaders on the Glue leaderboard for CoLA and QQP tasks. On two additional tasks (MNLI-Mismatched and RTE), the new model was significantly better than BERT-base on which it is based. We hypothesize that semantic tokenization has significant advantages in single-sentence or sentence similarity tasks.

Two additional tasks showed poorer performance (SST-2 and WLNI). Analysis of related datasets showed very long sequences. Our hypothesis is that the small max sequence length used in the BERT-base training has a significant disadvantage in processing long sequences. Considering the limitation of the model (small model, small max sequence length), this represents significant advantages of semantic tokenization.

Benchmark	BERT-base	BERT+Semantic	Leader
CoLA	52.1	77.9	74.4
SST-2	93.5	88.1	97.8
MRPC	88.9/84.5	85.8 / 84.3	93.9/91.8
QQP	71.2	93.0 / 95.6	75.2/90.9
MLNI-m	84.6	82.0	91.9
MLNI-mm	84.6	88.0	91.4
QNLI	90.5	90.0	97.3
RTE	65.7	86.8	93.2
WLNI	56.3	56.3	95.9

Table 3: Glue Benchmark Results,

Word	BERT/WP	BERT/Semantic
condition	condition	condit, ion
conditions	conditions	condit, ions
conditioning	conditioning	condit, ioning
conditioned	conditioned	condit, ioned
conditional	conditional	condit, ional
conditioner	condition, er	condit, ioner
conditionality	conditional, ity	condit, ionality
conditionable	condition, able	condit, ionable
conditionally	condition, ally	condit, ionally

Table 4: Tokenization of Test Wordforms

6 Analysis of Tokenizer Performance

As shown above, the semantic tokenizer improved subword regularization efficiency and performed better on benchmarks. However, the root cause of this improvement was not clear. We devised two experiments to better understand the performance of the semantic tokenizer. First, We analyzed the quality of input embeddings generated by the semantic tokenizer with pretrained BERT WordPiece tokenizer. Next we evaluated the output embeddings generated by the BERT-base with semantic tokenizer and pretrained BERT.

6.1 Quality of Input Embeddings

Input embeddings of different forms of a root word should be similar to ensure semantics are represented efficiently. We selected 9 forms of the word condition (Table 4) and compared the new tokenizer embeddings with WordPiece/BERT. Figure 2 shows cosine similarities between embeddings across wordforms for each tokenizer. If a word was divided into subwords, we used mean pooling to compute its embedding. The cells are color coded to communicate the strength of similarity, with darker shading indicating higher similarity.

The new semantic tokenizer improved embedding similarity across different wordforms, despite using 40% more tokens. On average, the semantic tokenizer improved embeddings by 40% compared to WordPiece/BERT. Overall, these results prove that semantic tokenizer improves embeddings despite the increase in the number of tokens.

One explanation for this improvement is semantic concentration achieved through new subwords. The stem *condit* occurs more than 300,000 times in the Wiki corpus. However the word *conditioner* occurs only 876 time. Aggregation of stems and suffixes improves masking probability and hence embedding quality.

6.2 Sentence Semantics

We evaluated the impact of semantic tokenization on sentence level output embeddings. We composed 9 sentences with similar semantics using the 9 wordforms (Table 4). We compared output embeddings of the semantic tokenizer with WordPiece/BERT-base in Figure 3. We computed the output embedding by mean pooling tokens at the output layer for each model. On average, cosine similarities across sentences improved 11% compared to pretrained BERT-base/Wordpiece. Overall, these results prove that stemming improves the accuracy of sentence level embeddings.

7 Conclusion

In this paper, we demonstrate significant NLP performance through improved vocabulary formation and without increasing model size. Specifically, we develop a new tokenizer that adds semantics as a constraint to vocabulary formulation. The tokenizer is implemented as a drop-in replacement for the SentencePiece tokenizer. The tokenizer includes a trainer that uses stemming to improve subword formation. An encoder is developed that complements the semantic trainer and integrates semantic tokenization into NLP models. Overall, the new semantic tokenizer significantly improves NLP performance on processing, classification and search tasks without increasing model sizes. The new tokenizer more than doubles the number of wordforms represented in the vocabulary. The enhanced vocabulary significantly improves model convergence and quality of word embeddings. Our experimental using the tokenizer with BERT-base show top performance on two Glue tasks while being 1/20th of the size of the other leading models.

Acknowledgments

This work was performed in part under a US Air Force contract.

References

- [1] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter, 2020.
- [2] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Nee-lakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. *Advances in Neural Information Processing Systems* 33, 2020.
- [3] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2019, Minneapolis, MN, USA, June 2-7, 2019, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, 2019.
- [4] M. Brysbaert, Michaël A. Stevens, Paweł Mandera, and Emmanuel Keuleers. How many words do we know? practical estimates of vocabulary size dependent on word definition, the degree of language input and the participant’s age. *Frontiers in Psychology*, 7, 2016.
- [5] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. Deberta: Decoding-enhanced BERT with disentangled attention. *CoRR*, abs/2006.03654, 2020.
- [6] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural machine translation of rare words with subword units. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1715–1725, Berlin, Germany, August 2016. Association for Computational Linguistics.
- [7] Mike Schuster and Kaisuke Nakajima. Japanese and korean voice search. In *2012 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 5149–5152, 2012.
- [8] Taku Kudo. Subword regularization: Improving neural network translation models with multiple subword candidates. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 66–75, Melbourne, Australia, July 2018. Association for Computational Linguistics.
- [9] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. ALBERT: A lite BERT for self-supervised learning of language representations. In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net, 2020.
- [10] Taku Kudo and John Richardson. Sentencepiece: A simple and language independent subword tokenizer and detokenizer for neural text processing. In Eduardo Blanco and Wei Lu, editors, *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing, EMNLP 2018: System Demonstrations, Brussels, Belgium, October 31 - November 4, 2018*, pages 66–71. Association for Computational Linguistics, 2018.
- [11] Martin F Porter. Snowball: A language for stemming algorithms, 2001.
- [12] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, and Jamie Brew. Huggingface’s transformers: State-of-the-art natural language processing. *CoRR*, abs/1910.03771, 2019.

	Condition	Conditions	Conditioning	Conditioned	Conditional	Conditioner	Conditionality	Conditionable	Conditionally		
Condition	1.00	0.91	0.80	0.79	0.87	0.78	0.73	0.75	0.82		Semantic Tokenizer / BERT-base
Conditions		1.00	0.84	0.83	0.88	0.82	0.80	0.80	0.85		
Conditioning			1.00	0.93	0.86	0.88	0.86	0.87	0.87		
Conditioned				1.00	0.86	0.86	0.83	0.87	0.87		
Conditional					1.00	0.84	0.83	0.82	0.92		
Conditioner						1.00	0.86	0.87	0.86		
Conditionality							1.00	0.86	0.86		
Conditionable								1.00	0.85		
Conditionally									1.00		
Condition	1.00	0.67	0.43	0.48	0.41	0.79	0.43	0.79	0.41		WordPiece/BERT-Base
Conditions		1.00	0.41	0.47	0.44	0.60	0.47	0.61	0.44		
Conditioning			1.00	0.70	0.45	0.39	0.50	0.45	0.41		
Conditioned				1.00	0.55	0.46	0.55	0.54	0.50		
Conditional					1.00	0.45	0.84	0.47	0.84		
Conditioner						1.00	0.48	0.73	0.50		
Conditionality							1.00	0.51	0.80		
Conditionable								1.00	0.51		
Conditionally									1.00		

Figure 2: Wordform Embedding Similarities.

	Circumstances added a condition to his actions.	Circumstances added conditions to his	Circumstances were conditioning his actions.	Circumstances conditioned his actions.	Circumstances caused his action to be	Circumstances were a conditioner to his	Circumstances added conditionality to his	Circumstances made his actions conditionable.	Circumstances added conditionality to his		
Circumstances added a condition to his actions.	1.00	0.95	0.83	0.78	0.81	0.92	0.93	0.82	0.74		Semantic Tokenizer / BERT-base
Circumstances added conditions to his		1.00	0.89	0.86	0.84	0.94	0.97	0.89	0.80		
Circumstances were conditioning his actions.			1.00	0.93	0.79	0.89	0.88	0.93	0.74		
Circumstances conditioned his actions.				1.00	0.78	0.83	0.85	0.93	0.73		
Circumstances caused his action to be					1.00	0.81	0.85	0.84	0.73		
Circumstances were a conditioner to his						1.00	0.93	0.87	0.77		
Circumstances added conditionality to his							1.00	0.88	0.83		
Circumstances made his actions conditionable.								1.00	0.77		
Circumstances added conditionality to his									1.00		
Circumstances added a condition to his actions.	1.00	0.92	0.65	0.53	0.86	0.88	0.87	0.87	0.83		BERT Base
Circumstances added conditions to his		1.00	0.75	0.65	0.85	0.81	0.84	0.91	0.84		
Circumstances were conditioning his actions.			1.00	0.88	0.63	0.65	0.62	0.68	0.60		
Circumstances conditioned his actions.				1.00	0.55	0.46	0.55	0.54	0.50		
Circumstances caused his action to be					1.00	0.78	0.90	0.90	0.87		
Circumstances were a conditioner to his						1.00	0.79	0.77	0.72		
Circumstances added conditionality to his							1.00	0.83	0.95		
Circumstances made his actions conditionable.								1.00	0.85		
Circumstances added conditionality to his									1.00		

Figure 3: Sentence Embedding Similarity.